

Scale-Aware Evolutionary Discovery of Wind-Farm Layout Optimizers with Large Language Models

Author One Author Two Author Three

ADDRESS

Correspondence: (EMAIL)

Abstract

We present a framework for the automated discovery of *deployable* learning-rate and constraint-penalty schedules for gradient-based wind-farm layout optimization using large language models (LLMs). Candidate schedules are expressed in a *scale-aware* parameterization in which the exploration learning rate is constructed from the rotor diameter and the only user-supplied scale is the constraint tolerance, so a single schedule generalizes across farms of different size with no free per-farm parameter. An island-based, quality-diversity evolutionary loop drives LLM coding agents to propose schedules, which a multi-stage *cascade* evaluator scores across a multi-farm, multi-scale training suite under a hard feasibility gate. A behavioral archive, content-addressed caching, and complete lineage provenance make every discovery auditable and, on a fixed platform, bit-reproducible; a firewall discipline keeps held-out and test performance out of the search. We describe the schedule interface, the fixed optimization skeleton, the evolutionary loop and cascade evaluator, the fitness, selection, and deployment rules, and the reproducibility and firewall machinery that together turn opportunistic program search into a repeatable engineering procedure.

1 Introduction

When designing a wind farm, the placement of turbines within the buildable area determines every downstream stage of a project, from land valuation to annual energy production (AEP). Because a fraction-of-a-percent gain in modeled AEP can translate into a large change in a project’s financial case, layout optimization is a well-studied problem: it is a non-convex, non-linear, constrained optimization in which turbine positions interact through wake losses and must respect spacing and boundary (inclusion/exclusion zone) constraints.

Gradient-based solvers have become a de-facto method for this problem. The TopFarm stochastic gradient descent (SGD) optimizer (Quick et al., 2022) adapts the Adam algorithm (Kingma and Ba, 2015) with a scheduled learning rate and a scheduled linear constraint penalty. The *schedule* — how the learning rate, the penalty weight, and the two Adam moment-decay coefficients evolve over the course of the optimization — is a rich design space that governs the trade-off between exploring the layout and driving the final iterate onto the feasible set. It is precisely the kind of low-dimensional, high-leverage function that is difficult to design by hand and attractive to search automatically.

Large language models have recently been used as the mutation operator in evolutionary *program search*: FunSearch (Romera-Paredes et al., 2024) and AlphaEvolve (Google DeepMind, 2025) evolve populations of programs by repeatedly prompting an LLM to improve high-scoring candidates. We adopt this paradigm for optimizer-schedule design, and add three ingredients that make the discovered artifact *deployable* rather than merely high-scoring on one farm:

1. a **scale-aware schedule interface** (Section 2.2) that removes the free learning-rate parameter and builds the exploration scale from the rotor diameter, so a discovered schedule is scale-covariant and transfers across farms of different turbine size;
2. a **quality-diversity evolutionary loop with a cascade evaluator** (Sections 2.5–2.7) that scores candidates across a multi-farm, multi-scale training suite under a hard feasibility gate, rejecting infeasible or dominated schedules cheaply before spending on expensive high-count evaluations; and

41 3. a **provenance, firewall, and reproducibility layer** (Sections 2.9–5) — full lineage, a behavioral
 42 archive, content-addressed caching, a scoped mutator workspace, and a pre-registered selection
 43 and deployment criterion — so that any “the system discovered X ” claim is auditable against
 44 what was seeded and what was measured.

45 The remainder of the paper describes the framework (Section 2), the wind farms and evaluation
 46 cells it is applied to (Section 3), the schedules the search discovers (Section 4), and the reproducibility
 47 and cost-control machinery (Section 5).

48 2 Framework

49 This section describes the discovery framework: the optimization problem (2.1), the scale-aware
 50 schedule interface (2.2), the fixed optimization skeleton (2.3), the penalty normalization (2.4), the
 51 evolutionary loop (2.5) and its behavioral archive (2.6), the cascade evaluator (2.7), the fitness and
 52 feasibility rules (2.8), and the provenance, firewall, selection, and deployment machinery (2.9–2.11).
 53 Figure 1 gives the overall data flow.

54 2.1 Optimization problem

55 For a farm with N turbines at horizontal positions $(\mathbf{x}, \mathbf{y})$ we maximize the modeled annual energy
 56 production

$$\max_{\mathbf{x}, \mathbf{y}} J(\mathbf{x}, \mathbf{y}), \quad J = \frac{8760}{10^6} \sum_k w_k P(\mathbf{x}, \mathbf{y}; u_\infty^{(k)}, \theta^{(k)}), \quad (1)$$

57 where P is the wind-farm power under inflow speed $u_\infty^{(k)}$ and direction $\theta^{(k)}$ with probability weight
 58 w_k , evaluated with an engineering wake model (Bastankhah and Porté-Agel, 2014), subject to the
 59 feasibility constraints

$$d(x_i, y_i) \geq 0 \quad \forall i, \quad \sqrt{(x_i - x_j)^2 + (y_i - y_j)^2} \geq D_{\min} \quad \forall i \neq j, \quad (2)$$

60 where $d(x_i, y_i)$ is the signed distance from turbine i to the nearest inclusion-zone boundary (negative
 61 outside the buildable region) and $D_{\min}$ is the minimum inter-turbine spacing. Feasibility is judged
 62 against a tolerance $\gamma_{\min}$ (in metres): a layout is feasible when every turbine satisfies the boundary
 63 constraint to within $\gamma_{\min}$ and the pairwise spacing constraint. $\gamma_{\min}$ is the single user-supplied scale of
 64 the problem.

65 2.2 Scale-aware schedule interface

66 The object we search over is a pure function that returns, at each optimization step i , the four
 67 quantities the skeleton needs: a learning rate η_i , a constraint penalty multiplier α_i , and the two Adam
 68 moment-decay coefficients $\beta_{1,i}, \beta_{2,i}$:

$$(\eta_i, \alpha_i, \beta_{1,i}, \beta_{2,i}) = S(i, T, D, D_{\min}, N, \gamma_{\min}, \alpha_0). \quad (3)$$

69 The arguments are the current step i , the total number of steps T , and the *problem-intrinsic scales*:
 70 the rotor diameter D , the minimum spacing $D_{\min}$, the turbine count N , the metre-valued constraint
 71 tolerance $\gamma_{\min}$, and a skeleton-computed penalty normalizer α_0 (Section 2.4). The Python signature is
 72 reproduced in Appendix A.

73 Two design choices make a discovered schedule deployable across scales:

Algorithm 1 Fixed optimization skeleton (the search controls only $\mathcal{S}$)

```
1:  $s \leftarrow \text{wind-aware init}(D, D_{\min}, N, \text{resource}); \quad \mathbf{m}, \mathbf{v} \leftarrow \mathbf{0}$ 
2:  $\alpha_0 \leftarrow \text{mean}(|\nabla_s J(s)|)/D; \quad \alpha_0 \leftarrow \text{CANONICALIZE}(\alpha_0)$  ▷ Section 2.4
3: for  $i = 0, 1, \dots, T - 1$  do
4:    $(\eta_i, \alpha_i, \beta_{1,i}, \beta_{2,i}) \leftarrow \mathcal{S}(i, T, D, D_{\min}, N, \gamma_{\min}, \alpha_0)$ 
5:    $\mathbf{j} \leftarrow \nabla_s [-J(s)] + \alpha_i \nabla_s \gamma(s)$  ▷ objective + penalty gradient
6:    $\mathbf{m} \leftarrow \beta_{1,i} \mathbf{m} + (1 - \beta_{1,i}) \mathbf{j}; \quad \mathbf{v} \leftarrow \beta_{2,i} \mathbf{v} + (1 - \beta_{2,i}) \mathbf{j}^2$ 
7:    $\hat{\mathbf{m}} \leftarrow \mathbf{m} / (1 - \beta_{1,i}^{i+1}); \quad \hat{\mathbf{v}} \leftarrow \mathbf{v} / (1 - \beta_{2,i}^{i+1})$ 
8:    $s \leftarrow s - \eta_i \hat{\mathbf{m}} / \sqrt{\hat{\mathbf{v}}}$ 
9: end for
10: return  $s$ 
```

No free learning rate. There is no driver-supplied initial learning rate η_0 anywhere in Equation (3). The exploration scale must be constructed *inside* the schedule from the rotor diameter, e.g. a peak learning rate $\eta_{\text{peak}} = c D$ with c a schedule-internal constant or curve that the search may evolve but that is never a user knob. Because the turbine spacing, wake length, and feasible-region geometry all scale with D , expressing the exploration scale as a multiple of D makes the schedule scale-covariant: the same function applied to a farm with a different rotor diameter automatically takes appropriately sized steps.

Tolerance as the terminal learning rate. Following TopFarm’s own convention, the terminal learning rate is a proxy for how tightly the boundary is ultimately satisfied; we expose it directly as the metre-valued tolerance $\gamma_{\min}$ toward which the schedule decays the learning rate. It is the only user-supplied scale, and a faithful schedule must respond to it: at a looser $\gamma_{\min}$ the schedule should stop decaying sooner, and its terminal behavior and feasibility should change accordingly (this responsiveness is checked in Section 2.7).

2.3 Fixed optimization skeleton

The schedule is the *only* thing the search controls. A fixed skeleton (Algorithm 1) handles initialization, gradient computation, and the Adam update. It (i) places the initial layout on a wind-direction-aware grid inside the buildable region (for the multi-zone case, on a zone-interior fill); (ii) forms the objective gradient of Equation (1) and the constraint gradient of the boundary and spacing penalties by automatic differentiation; (iii) combines them as $\mathbf{j} = \nabla_{\text{obj}} + \alpha_i \nabla_{\text{con}}$; and (iv) applies an Adam step with the scheduled $\eta_i, \beta_{1,i}, \beta_{2,i}$. The horizon is fixed at $T = 8000$ steps. Nothing in the skeleton contains a hardcoded learning rate; the entire step-size behavior comes from Equation (3).

2.4 Penalty normalization

So that the penalty multiplier is on a problem-intrinsic scale, the skeleton computes the normalizer

$$\alpha_0 = \frac{\text{mean}(|\nabla_x J|, |\nabla_y J|)}{D} \quad (4)$$

at the initial layout and passes it into Equation (3); the schedule composes the running penalty α_i from α_0 . Normalizing the gradient magnitude by the rotor diameter (rather than by a learning rate) keeps the penalty scale covariant with the geometry and independent of the step-size choice, consistent with the removal of a free learning rate. To make evaluation bit-reproducible across *independent processes on a given platform*, α_0 is canonicalized to single precision at the skeleton boundary (CANONICALIZE in Algorithm 1): the double-precision gradient reduction agrees to far more significant figures than single precision, so rounding α_0 to a canonical single-precision value makes the same (schedule, cell, seed)

Algorithm 2 Evolutionary schedule discovery

```
1: seed each island archive with generation-0 ancestors (logged with lineage)
2: while cost < ceiling do
3:   for each island do
4:      $p \leftarrow \text{SAMPLEPARENT}(\text{archive})$  ▷ fitness/novelty-aware
5:     materialize a scoped mutator workspace (Section 2.10)
6:      $c \leftarrow \text{LLM-MUTATE}(p, \text{firewall-safe feedback})$ 
7:     if NOVELTYFILTER rejects  $c$  then continue
8:     end if
9:      $(\text{fit}, \text{feasible}, \text{desc}) \leftarrow \text{CASCADE}(c)$  ▷ Algorithm/Section 2.7
10:    if feasible then ARCHIVEINSERT(island,  $c$ , fit, desc)
11:    end if
12:    LOGLINEAGE( $c$ ,  $p$ , engine, tokens, cost, desc, fit)
13:  end for
14:  periodically migrate elites between islands; checkpoint state
15: end while
```

104 reproduce identically in independent processes. This matters because the low-learning-rate tail of the
105 optimization is sensitive to α_0 . Across *different* hardware architectures a small floating-point drift
106 remains; we therefore report cross-platform agreement to within a measured drift tolerance (Section 5)
107 rather than bit-identity.

108 2.5 Evolutionary discovery loop

109 Schedules are discovered by a quality-diversity evolutionary loop (Algorithm 2), built on an island
110 model with a MAP-Elites archive (Mouret and Clune, 2015) per island. Each iteration: (i) a parent
111 schedule is sampled from an island’s archive with a fitness- and novelty-aware sampler; (ii) an LLM
112 mutation engine is prompted with the parent source and firewall-safe feedback and returns an improved
113 schedule as Python source; (iii) a code-novelty filter rejects a child that is a near-duplicate of an
114 already-seen program before any expensive evaluation is spent; (iv) the survivor is scored by the
115 cascade evaluator (Section 2.7); and (v) it is inserted into the island’s behavioral archive if it is elite
116 for its cell. Islands periodically exchange elites. The loop runs until a token/cost ceiling is reached
117 (Section 5).

118 The population is *seeded* with known-good schedules, logged as generation-0 ancestors with their
119 port transforms (Section 2.9); the search is an engineering procedure and seeding is encouraged. The
120 mutation operators are LLM coding agents accessed through a subscription/OAuth path (the Claude
121 Agent SDK is the primary engine; a command-line Gemini engine and a deterministic mock engine
122 are also supported). We adopt two eval-saving mechanisms from ShinkaEvolve (Sakana AI, 2025):
123 code-novelty rejection-sampling and fitness/novelty-aware parent sampling; the island/MAP-Elites
124 and cascade scaffolding follow OpenEvolve (OpenEvolve, 2025).

125 2.6 Behavioral archive

126 Each candidate is placed in a MAP-Elites archive keyed by four behavioral descriptors computed from
127 the schedule function *before* evaluation, so that the archive covers qualitatively distinct schedule shapes
128 rather than crowding a single family (Table 1): the peak learning rate relative to the rotor diameter,
129 the terminal learning rate in metres, the coupling between the penalty and the learning rate, and the
130 number of learning-rate restarts. Within each island, one elite is kept per occupied cell; feasibility
131 dominates, then mean fitness, then a worst-cell tiebreak (Section 2.8). Seeding populates several
132 descriptor cells from generation 0.

Table 1: Behavioral descriptors and their (frozen) bins for the MAP-Elites archive. Descriptors are computed from the schedule function prior to evaluation.

Descriptor	Bins
Peak learning rate / D	< 0.5 , $0.5-0.8$, $0.8-1.2$, >1.2
Terminal learning rate (m)	≤ 0.01 , $0.01-0.1$, $0.1-1$, >1
Penalty coupling	coupled, decoupled, cyclic
Learning-rate restarts	0, 1-2, ≥ 3

2.7 Cascade evaluator

Evaluating a schedule on every training cell at full seed count is expensive, so candidates pass through a cascade that spends compute in proportion to promise. Each stage runs the fixed skeleton (Algorithm 1) on a cell/seed and records the resulting AEP and feasibility.

Stage A (gross fast-reject). Two cheap cells $\times$ two seeds. This is a deliberately *gross* filter: a candidate is rejected only if any seed is infeasible or its AEP is more than about one percent below that cell’s reference. It is *not* texture-floor-tight — a floor-tight Stage A would mass-reject the very exploration the quality-diversity archive exists to keep, so the texture floors are reserved for selection margins, not for this filter. The rejection rate by cause (infeasible vs. below-reference) is logged as a pilot diagnostic. No high-count cell is used here.

Stage B (training matrix). The full frozen training set (Table 2) $\times$ at least five *paired* seeds — the same seed produces the same wind-aware initial layout for the candidate and for the reference, so per-seed differences are paired. Each cell yields a percentage score (Section 2.8) and a candidate-feasibility flag; the cross-cell aggregate is the candidate’s stage-B fitness.

Stage B⁺ (elite tier). Top- k archive elites are re-scored on expensive high-count cells, where a longer per-evaluation budget is affordable. This stage runs only on dedicated compute and never inside the main loop. It also carries the *capability-frontier* cells — cells (a high-count single-polygon farm and a dense multi-zone farm under a unidirectional resource) on which even the reference schedule is infeasible. These are tracked at the elite tier as qualitative probes: a candidate that reaches strict feasibility there is a qualitatively new result. They never gate the per-generation search; off the dedicated compute they are simply marked pending and deferred.

Stage C (holdout & responsiveness). Elites are scored on the held-out farm for margin-aware selection (Section 2.11) and are re-scored at a looser tolerance to confirm the schedule *responds* to $\gamma_{\min}$ (feasibility and terminal behavior must change); a schedule invariant to the tolerance is rejected regardless of AEP. The held-out AEP never leaves this stage: only feasibility booleans and a margin-over-floor boolean cross the firewall (Section 2.10).

2.8 Fitness and feasibility

For a training cell c , over the shared paired seed set, the score is the percentage improvement in mean AEP over a reference schedule,

$$\text{score}_c = 100 \times \frac{\overline{\text{AEP}}_{\text{cand},c} - \overline{\text{AEP}}_{\text{ref},c}}{\overline{\text{AEP}}_{\text{ref},c}}, \quad (5)$$

where the reference is the native TopFarm monotone schedule at a rotor-diameter-scaled learning rate, $\eta_{\text{peak}} = c_0 D$ with c_0 calibrated once on the training farm, decaying to $\gamma_{\min}$. Expressing fitness as percent-over-reference makes cells of very different magnitude (thousands of GWh for the large offshore farms, hundreds for the small onshore multi-zone farm) commensurable. The reference AEP

166 $\overline{\text{AEP}}_{\text{ref},c}$ is used purely as a *scale constant*: it normalizes the candidate even on a cell where the
 167 reference itself is infeasible, and the reference’s feasibility never confers credit.

168 Feasibility at $\gamma_{\min}$ is a *hard gate*: a candidate must be feasible on every paired seed of every training
 169 cell, or it is rejected (its score is set to $-\infty$ and it is excluded from the archive as infeasible). The
 170 cross-cell aggregate is *farm-balanced*: the mean over farms of each farm’s mean cell score, with a
 171 worst-cell tiebreak,

$$\text{fitness} = \frac{1}{|\mathcal{F}|} \sum_{f \in \mathcal{F}} \frac{1}{|C_f|} \sum_{c \in C_f} \text{score}_c, \quad \text{tiebreak} = \min_c \text{score}_c, \quad (6)$$

172 where $\mathcal{F}$ is the set of farms and C_f the scored cells on farm f . Averaging over farms rather than over
 173 cells makes each farm contribute equally irrespective of how many cells it carries (the training set has
 174 more cells on the large offshore farm than on the small multi-zone farm), so a percent improvement
 175 is worth the same on either farm; the worst-cell tiebreak still prevents a large win on one easy cell
 176 from masking a regression on a hard one. The aggregate is taken over the *scored* cells. A cell whose
 177 objective is saturated — for example a unidirectional resource in which every constraint-feasible layout
 178 that separates the turbines reaches free-stream power, so all such layouts score identically — carries no
 179 AEP-improvement signal; it is retained as a *feasibility-only* gate (evaluated at two seeds, participating
 180 in the hard feasibility gate) but is excluded from the scored set.

181 2.9 Provenance and lineage

182 Every candidate is recorded in an append-only lineage log with a content hash, its parent identifier(s),
 183 the mutation engine and model string, the token count and monetary cost, its behavioral descriptors,
 184 its per-cell fitness, and its generation and island. Seeded ancestors are logged as generation 0 together
 185 with their port transforms. Because the seeds and their transforms are recorded, any later claim that
 186 the system *discovered* a schedule family is measured against what was seeded, and “discovered” is
 187 distinguishable from “recombined a seed.”

188 2.10 Firewall and scoped workspace

189 The search must not overfit to, or leak, the held-out and test farms. Two disciplines enforce this.
 190 First, *information firewall*: only firewall-safe feedback — per-cell percentage scores and feasibility
 191 booleans — ever reaches a mutation engine; held-out and test AEP values never enter a mutator’s
 192 context or transcript, and the pre-registration of the selection and test design (Section 2.11) lives
 193 outside the search entirely. Second, *workspace scoping*: before each mutation the engine is run inside
 194 a freshly materialized clean-room directory that contains only the schedule interface, a sanitized copy
 195 of the skeleton and seed schedules, and the firewall-safe feedback; the results, analysis, held-out state,
 196 and pre-registration are outside the engine’s working directory and tool scope. Sanitization strips
 197 provenance and redacts any residual path token from the reference code, and a gate check refuses to
 198 launch if a forbidden token or a syntactically invalid reference file is present in the scoped workspace.

199 2.11 Selection and deployment criterion

200 Selection is margin-aware and never a single-evaluation argmax. Among the archive elites that are
 201 feasible on the training and held-out farms, the deployed schedule is the one that maximizes the
 202 multi-seed validation-mean margin over the reference on the held-out farm, provided that margin
 203 exceeds the held-out cell’s evaluation-texture floor; if no elite clears the margin, the native seed
 204 is deployed and the run is recorded as producing no improvement. The deployment claim is then
 205 confirmed, once, on a pre-registered test set under a paired multi-seed protocol (identical per-seed
 206 initial layouts across arms) and, as a fidelity stage, re-validated in the full stochastic optimizer. Because
 207 feasibility is judged against a tolerance and reference results depend on it, all comparisons are reported
 208 at matched tolerance.




Figure 1: The discovery loop. Parents are sampled from per-island MAP-Elites archives; an LLM mutation engine, run inside a scoped clean-room workspace, proposes a child schedule from the parent source and firewall-safe feedback; a novelty filter rejects near-duplicates before evaluation; the cascade evaluator scores survivors; feasible elites are archived and every candidate is logged with full provenance. Held-out selection, the pre-registered paired test, and stochastic re-validation sit outside the loop.

3 Application

We apply the framework across three wind-farm geometries and four wind resources. The farms span a range of rotor diameters and constraint geometries: the Danish Energy Island (DEI) cluster with IEA 15 MW turbines ($D = 240$ m, single polygon); the IEA Wind 740-10-MW Reference Offshore Wind Plant (ROWP) with IEA 10 MW turbines ($D = 198$ m, single polygon), held out for selection; and the Parque Ficticio inclusion-zone case of Criado Risco et al. (2024) with Vestas V80 turbines ($D = 80$ m, five disconnected inclusion zones — the hardest feasibility geometry). Minimum spacing is four rotor diameters for DEI and ROWP and two rotor diameters for Parque. We consider four wind resources: the measured DEI and ROWP roses, an omnidirectional resource, and a unidirectional resource as limiting cases.

Training cells. The search optimizes and sees percentage feedback on the training cells of Table 2, which span two rotor diameters, turbine counts from 10 to 120, three resource types (the DEI rose, omnidirectional, and unidirectional — the ROWP rose is held out), the multi-zone geometry, and a high-count cell. Two cheap cells serve as the Stage-A fast-reject filter. The reference AEP column is the scale constant of Equation (5), and the floor column is the per-cell texture floor. One Parque cell under the unidirectional resource is a feasibility-only gate: its objective is saturated (Section 2.8), so it is scored for feasibility only and excluded from the aggregate. The two capability-frontier cells (Section 2.7) — a high-count DEI cell and a dense Parque cell under the unidirectional resource, on which even the reference is infeasible — are not training cells; they are elite-tier probes only.

Held-out and test farms. ROWP is an intermediate scale not present in training and is used only for margin-aware selection; its AEP is firewalled from the search. The pre-registered test set — touched exactly once, at the deployment decision — comprises the high-count ROWP cells ($N = 200$ and $N = 300$), the Parque site’s real heterogeneous wind resource, and a unidirectional extreme on the held-out ROWP farm.

Evaluation-texture floors. AEP is a mildly discontinuous function of the layout because the wake model masks each turbine’s contribution outside a cone; at millimetre-scale perturbations this yields a small texture in re-scored AEP. We measure this floor per farm and set the per-cell minimum meaningful margin accordingly: 0.3 GWh for the large offshore DEI farm, 0.1 GWh for the multi-zone

Table 2: Training cells (all references feasible on every seed). D in metres; Ref. AEP is the rotor-diameter-scaled reference (the scale constant of Equation (5)), a multi-seed feasible mean in GWh; Floor is the per-cell evaluation-texture floor (GWh). “feas.-only” marks the saturated-objective cell that gates feasibility but is excluded from the score aggregate.

Cell	Farm	N	D	Resource	Ref. AEP	Floor
DEI-50	DEI	50	240	DEI rose	5560.4	0.3
DEI-80	DEI	80	240	omnidirectional	8818.1	0.3
DEI-120	DEI	120	240	DEI rose	13030.0	0.3
DEI-50 uni.	DEI	50	240	unidirectional	5598.4	0.3
Parque-20	Parque	20	80	DEI rose	231.5	0.1
Parque-14 uni.	Parque	14	80	unidirectional	184.8 (<i>f.o.</i>)	0.1
Parque-10	Parque	10	80	omnidirectional	127.0	0.1

237 Parque farm, and 0.64 GWh for the held-out ROWP farm. Selection margins use these per-cell floors
 238 rather than a single global threshold; the gross Stage-A filter (Section 2.7) does not use them.

239 3.1 Reference multistart budget and its convergence

240 The score of Equation (5) compares a candidate against a single reference run per seed, but the baseline
 241 a *deployed* schedule must actually beat is stronger: the best feasible layout a practitioner would field
 242 from a *multistart* of K independent random initializations of the reference. The number of starts that
 243 defines a converged multistart is not arbitrary; we fix it per farm with the shuffle-saturation analysis of
 244 Quick et al. (2026). A pool of M feasible reference runs is collected on a farm; the pool is randomly
 245 reshuffled 10^3 times, and for each shuffle the running maximum AEP is recorded against the number
 246 of starts considered. The p -th percentile of that running maximum across shuffles is a conservative
 247 estimate of the best AEP obtained with that many starts — the first percentile is the value beaten 99%
 248 of the time — and the convergence budget K^* is the number of starts at which that percentile reaches
 249 within a tolerance of the pool maximum (made precise below). The pool is large enough to resolve K^*
 250 when $K^* \leq M/2$; otherwise the pool is doubled and the analysis repeated.

251 Because AEP is effectively continuous near the optimum, the single best run of a pool is an isolated
 252 point, so reaching the *exact* pool maximum saturates only at $\approx 0.99 M$ regardless of M . We therefore
 253 read saturation at a physically negligible tolerance ϵ (reaching within ϵ of the pool maximum) and
 254 report K^* across ϵ . The budget is then a property of the density of the near-optimum basin: if a
 255 fraction f_ϵ of runs land within ϵ of the best, the first-percentile budget approaches $\ln(0.01)/\ln(1 - f_\epsilon)$
 256 as the pool grows.

257 We calibrate on the DEI training farm and the held-out ROWP farm with pools of $M \approx 4000$ feasible
 258 reference runs each, collected as HPC array jobs (Figure 2 and Table 3). The per-pool convergence
 259 estimate is itself noisy at small pools — on the held-out farm the first-percentile budget is 420 ± 214
 260 at $M = 1024$ against ≈ 1086 at the deep pool — so it is the deep pools with 10^3 -shuffle averaging,
 261 not a single half-pool check, that make K^* trustworthy; at the deep pools the half-pool rule holds
 262 on both farms ($K_1^* \leq M/2$), confirming the pool is large enough to resolve the budget. ROWP is
 263 consistently the harder farm: its near-optimum basin is sparser (0.39% of runs within 0.05% of the
 264 best against 0.52% for DEI, and 3.5% vs. 9.3% within 0.1%), so it requires 1.3–3 $\times$ more starts than
 265 the training farm. At the tenth percentile the budget is ≈ 290 starts for DEI and ≈ 550 for ROWP; for
 266 first-percentile confidence within 0.05% of the optimum it is ≈ 520 and ≈ 1090 (asymptotically ≈ 880
 267 and ≈ 1180). The 400-start choice of Quick et al. (2026) for their DEI farm falls between our tenth-
 268 and first-percentile DEI budgets, consistent with this calibration.

269 We adopt a reference multistart of $K = 1024$ starts. This exceeds the first-percentile budget on
 270 the training farm and the fifth-percentile budget on the held-out farm ($K_5^* \approx 730$), and covers the
 271 held-out first percentile at the 0.1% tolerance (≈ 130); at the strictest 0.05% tolerance it sits just




Figure 2: Shuffle-saturation convergence of the reference multistart on the DEI training farm (left) and the held-out ROWP farm (right), following Quick et al. (2026). Faint lines are individual shuffles of the pool; the black curves are the tenth-, fifth-, and first-percentile running maxima; the shaded band is within $\epsilon = 0.05\%$ of the pool maximum, and the dotted vertical line is the half-pool budget $M/2$. Each percentile reaches the band well inside $M/2$ on both farms; the held-out ROWP farm (sparser near-optimum basin) requires more starts than the training farm.

Table 3: Reference-multistart convergence budget per farm. The feasible pool has $M \approx 4000$ runs (DEI budgets resolved on a $T=2048$ sub-pool, ROWP on the full $T=4096$ pool); “spread” is the relative AEP range (max – min)/max over the pool; f_ϵ is the fraction of runs within ϵ of the pool maximum; $K_p^\star$ is the pool-measured number of starts at which the p -th-percentile running maximum reaches within ϵ of the pool maximum. The held-out ROWP farm needs 1.3–3 $\times$ more starts than the DEI training farm.

Farm	spread (%)	$f_{0.05\%}$ (%)	$f_{0.1\%}$ (%)	$K_{10}^\star$ (0.05%)	$K_1^\star$ (0.05%)	$K_1^\star$ (0.1%)
DEI-50 (training)	0.282	0.52	9.3	290	520	47
ROWP (held-out)	0.456	0.39	3.5	546	1086	130

below the held-out first-percentile budget (≈ 1086), a shortfall we regard as immaterial at that margin. Deployment comparisons (Section 2.11) are made against the best-of- K of this multistart at matched start count, not against its mean.

4 Discovered schedules

Figure 3 shows a schedule search on the training farm for two mutation engines. Each point is one candidate the LLM proposed and ran through the fixed skeleton; filled points are feasible on every seed, hollow points are infeasible, and the solid line is the best feasible fitness so far. Most candidates land within a few hundredths of a percent of the native reference, the best-so-far improves by only ≈ 0.05 – 0.09% over hundreds of attempts, and the infeasible outliers near $+0.7\%$ are exactly the non-deployable layouts the feasibility gate (Section 2.8) discards — a reminder that raw AEP off the feasible set is not a fitness signal. The picture is a faithful near-null: the interface admits broad exploration, but the scale-aware native reference is already close to the best a schedule achieves on this farm.

The value of the search is less the marginal AEP than the *diversity of deployable schedule shapes* it surfaces. Figure 4 plots four schedules evaluated at the held-out ROWP scales: the native reference (constant learning rate then a compounding-product decay, penalty coupled as $\alpha \propto 1/\eta$, fixed TopFarm moment coefficients $\beta_1=0.1, \beta_2=0.2$); a portfolio-discovered schedule (it21) that ramps quickly to the




Figure 3: Schedule search on the training farm for two mutation engines (Claude, left; Antigravity/Gemini 3, right). Each point is a candidate schedule proposed by the LLM and scored by the fixed skeleton; filled points are feasible on every seed, hollow points are infeasible; the solid line is the best feasible fitness so far; the dashed line is the native reference (0%); the star marks the best (deployed) schedule. The best-so-far improves only $\approx 0.05\text{--}0.09\%$ over the native reference, and the infeasible $+0.7\%$ outliers are non-deployable layouts discarded by the feasibility gate.

289 exploration scale, holds, then decays smoothly while *raising* the adaptive coefficient β_2 from 0.2 to
 290 0.77 and easing the momentum β_1 down into the tail; and two seeded ancestors that mark the extremes
 291 the interface admits — a dual-bump learning rate (it192) and a cyclic warm-restart schedule (it118)
 292 that repeatedly re-heats the learning rate and both moment coefficients. All four decay the learning
 293 rate toward the tolerance $\gamma_{\min}$ and restore the penalty near the end to drive the final iterate onto the
 294 feasible set, but they reach it by very different routes. That distinct, hand-inscrutable shapes all attain
 295 comparable feasible AEP is itself a finding: the objective is forgiving in schedule space once the
 296 scale-aware exploration and the terminal feasibility restoration are right.

297 5 Reproducibility and cost control

298 The framework is engineered to be re-runnable and auditable. A content-addressed cache keyed by
 299 (schedule hash, cell, seed, $\gamma_{\min}$, T) returns a stored result for any previously computed evaluation, so a
 300 resumed run recomputes nothing. State — the archives, the novelty seen-set, the cost tally, and the
 301 generation counter — is checkpointed atomically each generation; the parent-sampling randomness is
 302 derived from the run identifier and generation, and parent lists are ordered canonically, so a resumed
 303 run reproduces the same candidates and hits the cache for every evaluation. The single-precision
 304 canonicalization of α_0 (Section 2.4) makes each evaluation *bit-identical across independent processes*
 305 *on a fixed platform*. Across different hardware architectures a small floating-point drift remains —
 306 of order ± 1.7 GWh on the large offshore farm between arm64 and x86 in our measurements — so
 307 cross-platform results are compared at that measured drift tolerance rather than as bit-identity. A cost
 308 ceiling tracks cumulative tokens and spend from the per-invocation engine logs and cleanly aborts the
 309 run near the ceiling (it stops issuing new mutations, finishes in-flight work, checkpoints, and stops).
 310 Together these make a discovery run a deterministic function of its seeds, configuration, and budget on a
 311 fixed platform.




Figure 4: Four schedules evaluated at the held-out ROWP scales ($D = 198$ m, $D_{\min} = 792$ m, $N = 74$, $\gamma_{\min} = 0.01$, and α_0 from the skeleton): learning rate η , penalty weight α (log scale), and the two Adam moment coefficients β_1, β_2 over optimization progress $t = i/T$. The native reference (dashed) is shown against a portfolio-discovered schedule (it21) and two seeded ancestors (dual-bump it192; cyclic warm-restart it118). The interface admits widely different shapes; all decay η toward $\gamma_{\min}$ and restore the penalty terminally.

6 Conclusions

We described a framework for discovering deployable wind-farm layout optimizers with large language models. Its distinguishing features are a scale-aware schedule interface that removes the free learning rate and builds the exploration scale from the rotor diameter; a quality-diversity evolutionary loop whose cascade evaluator scores candidates across a multi-farm, multi-scale training suite under a hard feasibility gate; and a provenance, firewall, and reproducibility layer — lineage, a behavioral archive, content-addressed caching, a scoped mutator workspace, and a pre-registered selection and deployment criterion — that make the procedure auditable and repeatable. The framework turns LLM-driven program search from an opportunistic experiment into an engineering procedure whose output is a schedule ready to be deployed and confirmed on unseen farms.

A Schedule interface and mutation prompt

The searched object is a single Python function; the skeleton, seeds, and firewall-safe feedback are provided read-only in the scoped workspace.

```
def schedule_fn(step, total_steps, D, min_spacing,
               n_turbines, gamma_min, alpha0):
    """Return (lr, alpha, beta1, beta2) at 'step'.
```

D : rotor diameter (m) -- exploration-scale source

```

330     min_spacing: minimum inter-turbine spacing (m)
331     n_turbines : N                                -- density / packing signal
332     gamma_min  : constraint tolerance (m)          -- the terminal learning rate,
333                                                    the only user-supplied scale
334     alpha0      : penalty normalizer = mean(|grad J|)/D (skeleton-computed)
335     """"
336     ...
337     return lr, alpha, beta1, beta2

```

338 The mutation engine receives, in its scoped workspace: the interface above, a sanitized copy of the fixed
339 skeleton and the seed schedules, the parent schedule source, and the firewall-safe feedback (per-cell
340 percentage scores and feasibility booleans). It returns one improved module defining `schedule_fn`.
341 No held-out or test AEP, and none of the analysis or pre-registration, is ever placed in that workspace.

342 References

- 343 Bastankhah, M. and Porté-Agel, F.: A new analytical model for wind-turbine wakes, *Renewable Energy*, 70,
344 116–123, 2014.
- 345 Criado Risco, J., Valotta Rodrigues, R., Friis-Møller, M., Quick, J., Mølgaard Pedersen, M., and Réthoré, P.-E.:
346 Gradient-based wind farm layout optimization with inclusion and exclusion zones, *Wind Energy Science*, 9,
347 585–600, 2024.
- 348 Google DeepMind: AlphaEvolve: a coding agent for scientific and algorithmic discovery, Technical report,
349 2025.
- 350 Kingma, D. P. and Ba, J.: Adam: a method for stochastic optimization, in *International Conference on Learning*
351 *Representations*, 2015.
- 352 Mouret, J.-B. and Clune, J.: Illuminating search spaces by mapping elites, arXiv preprint arXiv:1504.04909,
353 2015.
- 354 OpenEvolve: an open-source evolutionary coding framework, software, 2025.
- 355 Quick, J., Réthoré, P.-E., Mølgaard Pedersen, M., and Rodrigues, R. V.: Stochastic gradient descent for wind
356 farm optimization, *Wind Energy Science Discussions*, 1–17, 2022.
- 357 Quick, J., Réthoré, P.-E., Maak, L., and Maheshwari, P.: Wind farm design under uncertainty: managing
358 unknown neighboring farm effects, in *Proceedings of the ASME 2026 45th International Conference on*
359 *Ocean, Offshore and Arctic Engineering (OMAE)*, OMAE2026-176207, 2026.
- 360 Romera-Paredes, B., Barekatin, M., Novikov, A., Balog, M., Kumar, M. P., Dupont, E., Ruiz, F. J., Ellenberg, J.
361 S., Wang, P., Fawzi, O., et al.: Mathematical discoveries from program search with large language models,
362 *Nature*, 625, 468–475, 2024.
- 363 Sakana AI: ShinkaEvolve: sample-efficient evolutionary program search with large language models, soft-
364 ware/technical report, 2025.